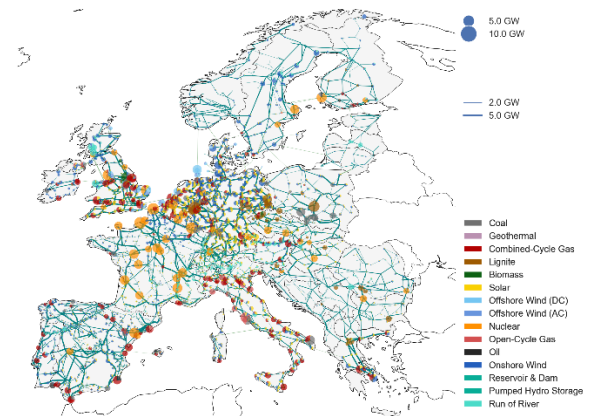


# Energy System Modeling with Python

University of Freiburg (Germany) | Faculty of Engineering  
Department of Sustainable Systems Engineering | INATECH  
**Chair for Control and Integration of Grids**

Tuesday, 3. June 2025



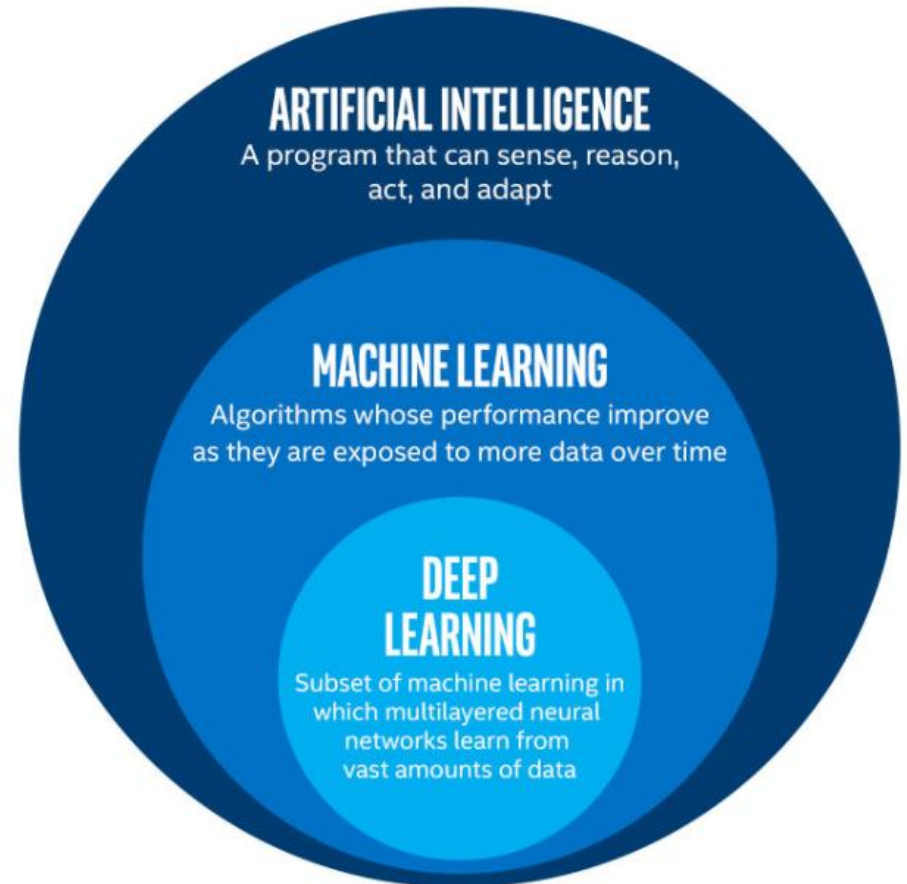
A stylized background graphic featuring a large, light blue leaf-like shape on the right and a cluster of overlapping hexagons on the left, all in shades of blue and teal.

# Introduction

# What is AI and ML?

---

- AI refers to the broader field of creating machines or software that can perform tasks that typically require human intelligence
- AI aims to mimic human cognitive functions in a way that allows machines to make autonomous decisions and perform tasks without explicit programming
- ML is a subfield of AI, which improves its performance on a task through experience (data)



# Let's define Machine Learning

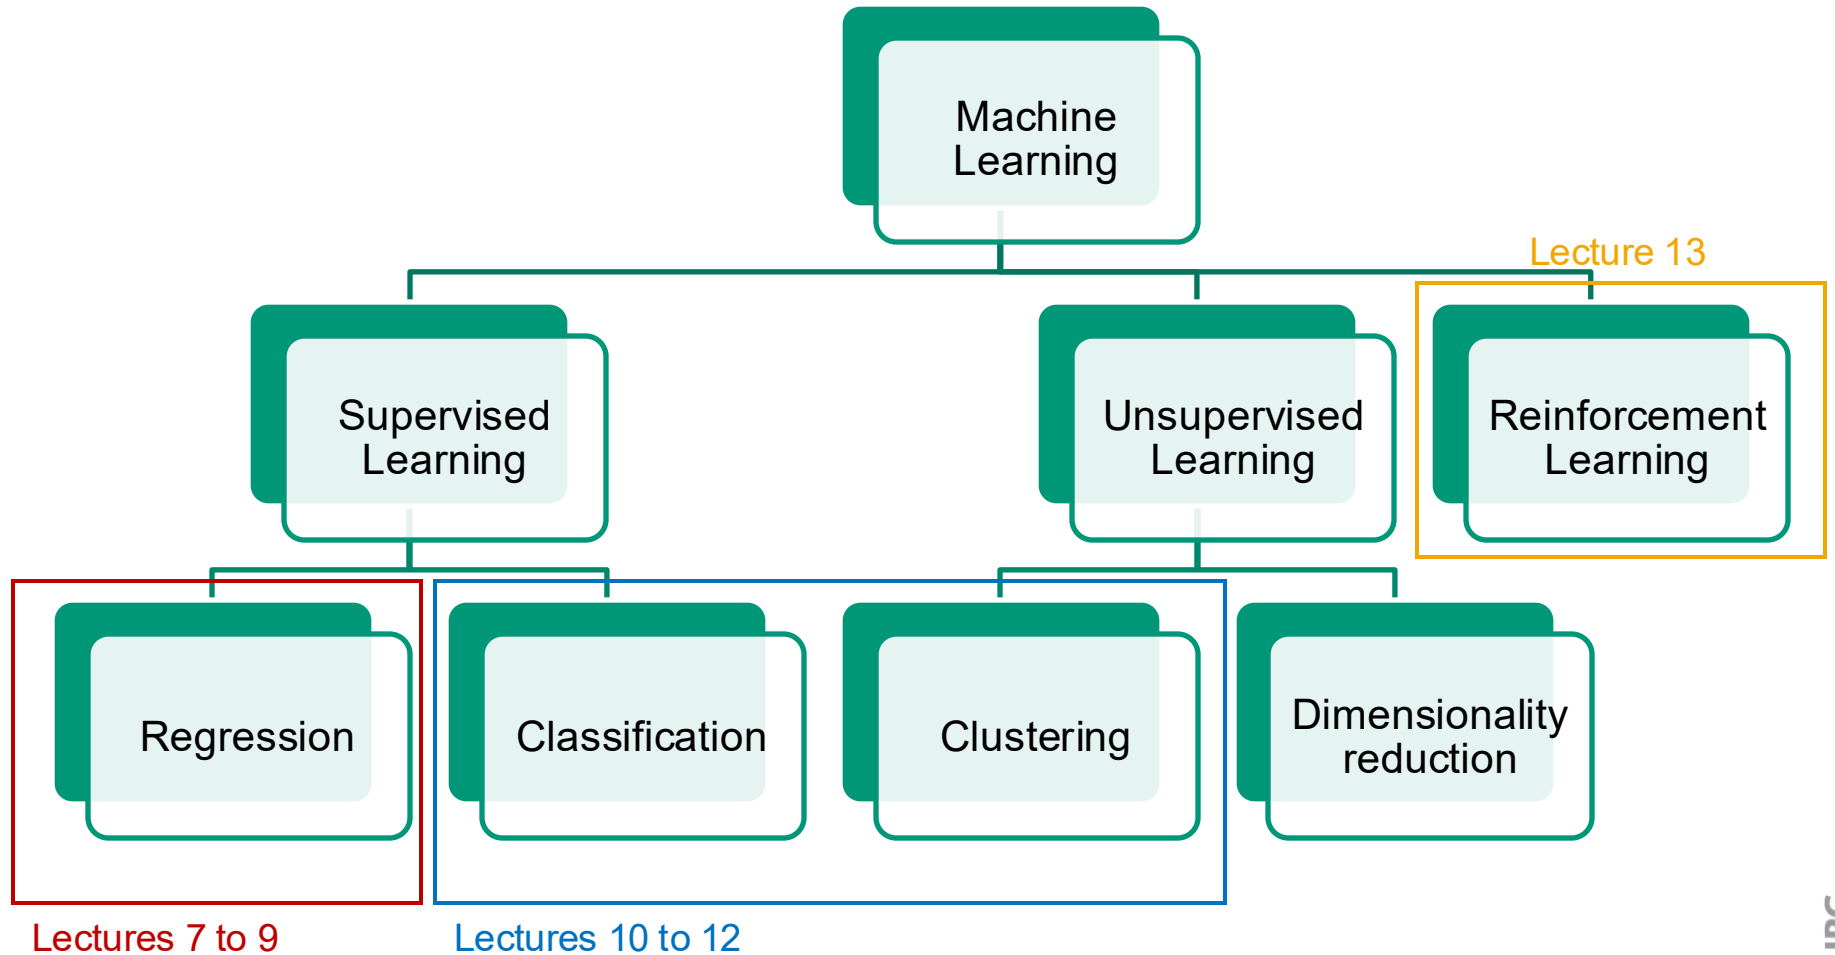
---

## *Tom Mitchell (1998):*

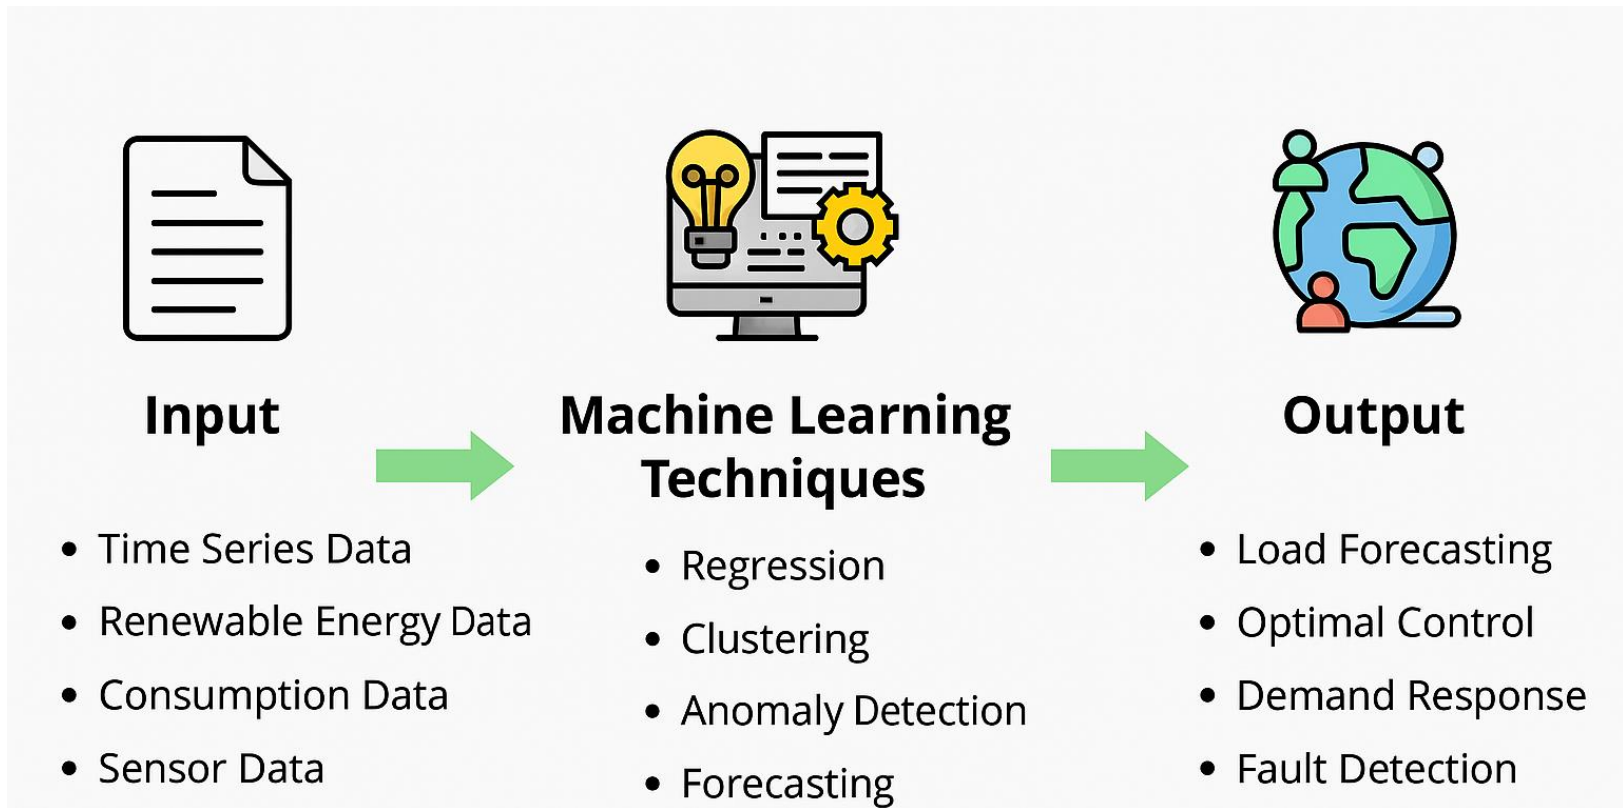
Machine learning is a well-posed learning problem where a computer program is said to **learn** from **experience E** with respect to some **task T** and some **performance measure P**, if its performance on T, as measured by P, improves with experience E.

So, what is **Experience E**, **Task T**, and **Performance measure P**?

# Branches of ML



# What is the process of ML?



So, what is **Experience E**, **Task T**, and **Performance measure P**?

# How can ML help us in the field of energy systems?

---

## *Energy Systems*

- Complex, nonlinear, and coupled dynamics
- High uncertainty from renewables and demand
- Incomplete or limited observational data
- Intractable to simulate all operating conditions
- Need for real-time control and decision-making

## *ML Capabilities*

- Learns complex behaviors without explicit system knowledge
- Adapts to variability and uncertainty
- Generalizes well from sparse or noisy data
- Approximates outcomes quickly without full models
- Enables fast, data-driven decisions in real time

# Limitations of ML in Energy Systems

## *Energy Systems*

- Energy systems are critical infrastructure
- Decisions must be explainable and auditable
- Operations must be reliable under all scenarios
- Failure can have serious physical and societal consequences
- Subject to strict regulations and accountability

## *ML Capabilities*

- ML models must be trustworthy, safe, and robust
- Many ML models, especially deep learning, are black boxes
- ML cannot guarantee behavior in rare or unseen cases
- ML errors may be hard to detect and correct in real time
- The EU AI Act classifies energy systems as high-risk applications

Before applying ML, **carefully assess whether it's the right tool for the job** — interpretability, robustness, and compliance must come first.



# Some terminology to remember

---

## Data and Learning Setup

- **Features:** Input variables used by the model to make predictions.
- **Labels:** True output values used in supervised learning.
- **Model:** A function that maps inputs (features) to predicted outputs.
- **Training:** Process where the model learns from data by minimizing error

## Model Evaluation

- **Testing / Evaluation:** Applying the trained model to unseen data to assess its generalization performance.
- **Performance Metrics:** Quantitative measures of model quality, such as:
  - **Accuracy** (classification)
  - **Precision / Recall / F1-score** (imbalanced classification)
  - **RMSE / MAE** (regression)
- **Loss Function:** A function that quantifies the difference between predicted and true values during training (e.g., Mean Squared Error, Cross-Entropy). The model learns by minimizing this.



Supervised  
Learning

***Regression***

# How supervised learning works

## Terminology:

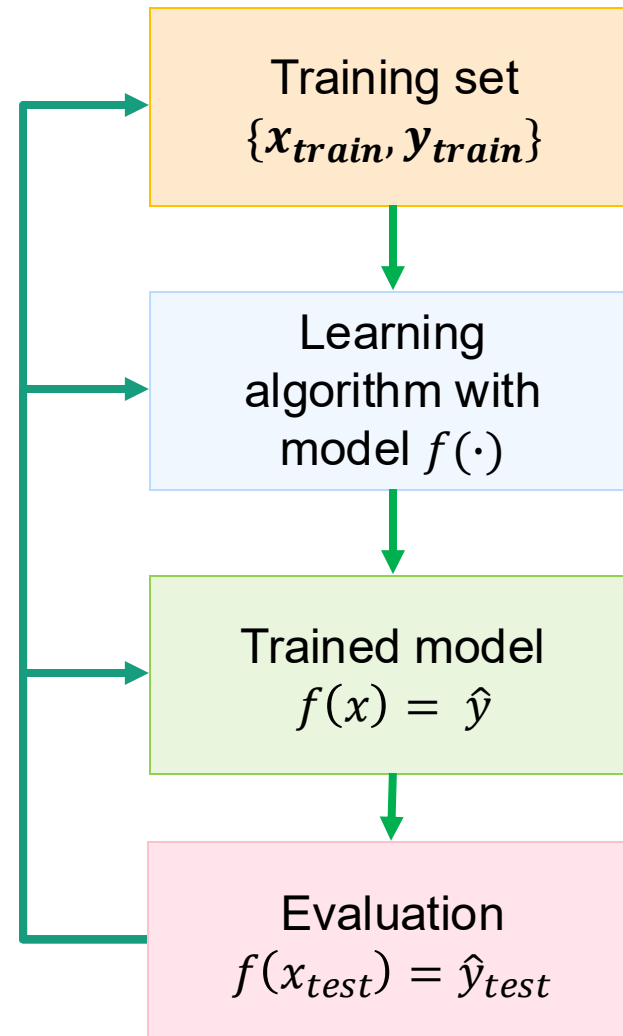
$X$ : features vector, where  $x \in X$

$Y$ : labels vector, where  $y \in Y$

$x_{train}, y_{train}$  : training dataset

$x_{test}, y_{test}$  : testing dataset that the model has never seen before

$f(\cdot)$  : a model used to approximate the values



# Supervised learning in simplest form: linear regression

---

The most basic equation to describe our model is known as **linear regression equation**:

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n = \beta_0 + \sum_{i=1}^n \beta_i x_i$$

where:

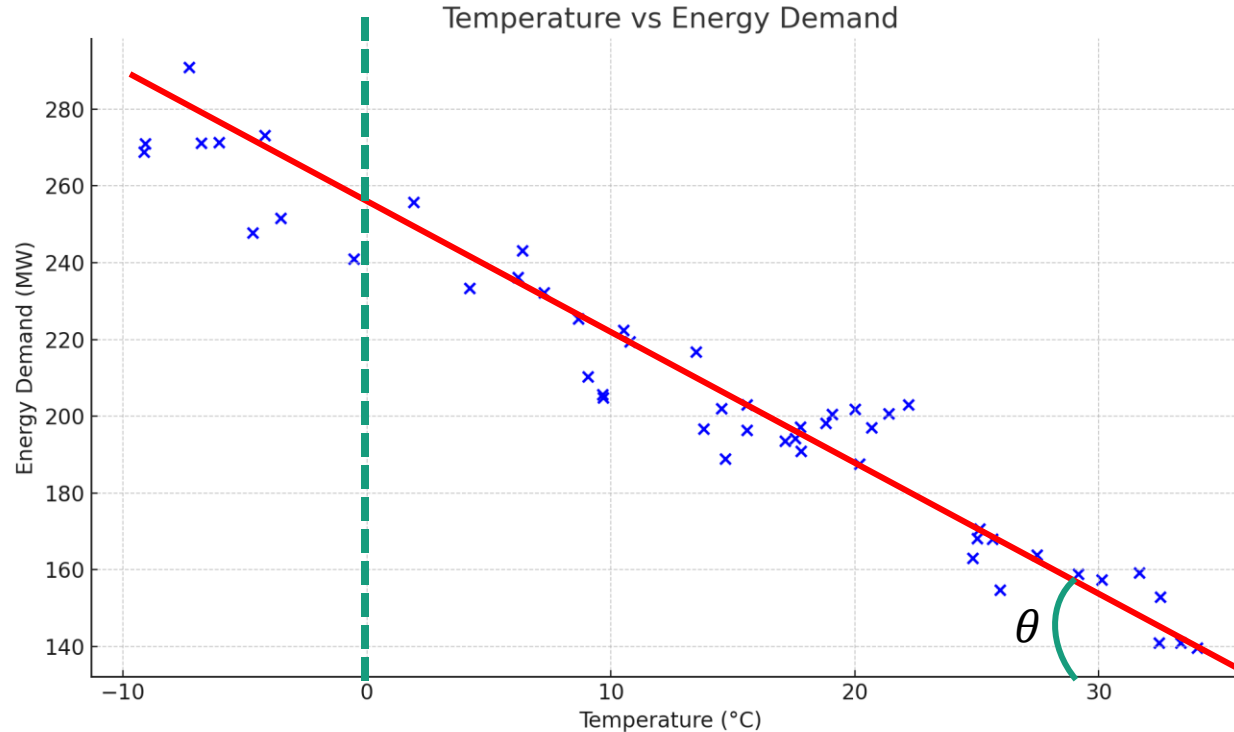
$\hat{y}$ : predicted **output** (e.g., energy demand),

$x_1, x_2, \dots, x_n$ : **input features** (e.g., temperature, wind speed, etc.),

$\beta_0$ : **bias term** (intercept),

$\beta_1, \beta_2, \dots, \beta_n$ : **coefficients** corresponding to each input feature.

# Supervised learning in simplest form: linear regression



$$\hat{y} = \beta_0 + \beta_1 x_1$$

$$\beta_0 = y \text{ when } x = 0$$

$$\beta_1 = \tan(\theta)$$

# How to train the model: finding the parameters

---

## Steps:

1) Define a model:  $\hat{y} = \beta_0 + \beta_1 x_1$

2) Define a loss function as **Squared Error**:

$$\mathcal{L}^{(i)} = (y^{(i)} - \hat{y}^{(i)})^2 = (y^{(i)} - (\beta_0 + \beta_1 x^{(i)}))^2$$

3) Define a cost function as **Mean Squared Error (MSE)**:

$$J(\beta_0, \beta_1) = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - \hat{y}^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - (\beta_0 + \beta_1 x^{(i)}))^2$$

4) Solve the equation:

- **Analytically**: Take derivatives of  $J$  with respect to  $\beta_0$  and  $\beta_1$ , set them to zero, and solve the resulting equations,
- **Numerically**: Use an iterative method like **gradient descent** to minimize the cost

# Minimizing the Cost — Gradient Descent

We want to minimize the *Mean Squared Error (MSE)*:

$$J(\beta_0, \beta_1) = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - \hat{y}^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - (\beta_0 + \beta_1 x^{(i)}))^2$$


We will apply **Gradient Descent**: Iteratively update parameters

$$\beta_j \leftarrow \beta_j - \alpha \cdot \frac{\partial J}{\partial \beta_j}$$

where:

$\alpha$ : is the **learning rate** (how fast we update our parameters)

**Gradient descent** is an iterative optimization algorithm that adjusts the model parameters step by step in the direction that **reduces the error** the most. It uses the **slope (gradient)** of the cost function to decide how to update the parameters, gradually moving toward the minimum.



# Minimizing the Cost — Gradient Descent

**Compute the partial derivatives:**

$$\frac{\partial J}{\partial \beta_0} = -\frac{2}{m} \sum_{i=1}^m \left( y^{(i)} - (\beta_0 + \beta_1 x^{(i)}) \right)$$
$$\frac{\partial J}{\partial \beta_1} = -\frac{2}{m} \sum_{i=1}^m \left( y^{(i)} - (\beta_0 + \beta_1 x^{(i)}) \right) \cdot x^{(i)}$$

**Gradient descent update rules:**

$$\beta_j \leftarrow \beta_j - \alpha \cdot \frac{\partial J}{\partial \beta_j}$$
$$\beta_0 \leftarrow \beta_0 + \frac{2\alpha}{m} \sum_{i=1}^m \left( y^{(i)} - (\beta_0 + \beta_1 x^{(i)}) \right)$$
$$\beta_1 \leftarrow \beta_1 + \frac{2\alpha}{m} \sum_{i=1}^m \left( y^{(i)} - (\beta_0 + \beta_1 x^{(i)}) \right) \cdot x^{(i)}$$



# How to evaluate the performance

## Mean Absolute Error (MAE)

### Definition:

MAE measures the **average magnitude of the errors**, ignoring their direction.

### Formula:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y^{(i)} - \hat{y}^{(i)}|$$

### Key Points:

- Measures **average absolute deviation**
- **Linear penalty** for all errors
- Same units as the target (e.g., €/MWh)

### Rule of thumb:

**MAE** for simple, balanced accuracy;  
**RMSE** when large errors must be minimized.

## Root Mean Squared Error (RMSE)

### Definition:

RMSE is the square root of the **mean of squared errors**.

### Formula:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2}$$

### Key Points:

- **Quadratic penalty** — large errors weigh more
- Always  $\geq$  MAE
- Same units as the target (e.g., €/MWh)

# Coffee Break

---



# Time to put everything into code

**START  
CODING  
SESSION**



# What You'll Do in Code Today

---

## 1. Implement Linear Regression from Scratch

Understand how a model learns the best-fitting line using gradient descent

## 2. Visualize the Learning Process

Plot how the model gradually improves over time — see the cost drop and the line converge

## 3. Apply Regression to Real Energy Data

Use real-world demand and price data from the German day-ahead market

## 4. Compare Your Model to Scikit-learn

Benchmark your results against a standard ML library

## 5. Evaluate Forecast Quality

Calculate error metrics (MAE, RMSE) to assess how good your predictions are

## 6. Multiple linear regression

You will implement multiple regression with more features of your choice

# Takeaways

---

- ✓ You now understand what **supervised learning** is — and how it applies to energy systems.
  - ✓ You've seen how a **linear regression model** can forecast energy demand or prices from input features.
  - ✓ You implemented **gradient descent** — the key algorithm for training many ML models.
  - ✓ You explored how **learning rate and normalization** affect model convergence.
- 
- How do MAE and RMSE differ — and when would you use each?
  - How does gradient descent *know* when it's done learning?
  - In which real-world energy tasks might **simple models** like linear regression already be “good enough”?
  - And where might they fail — what can't they capture?