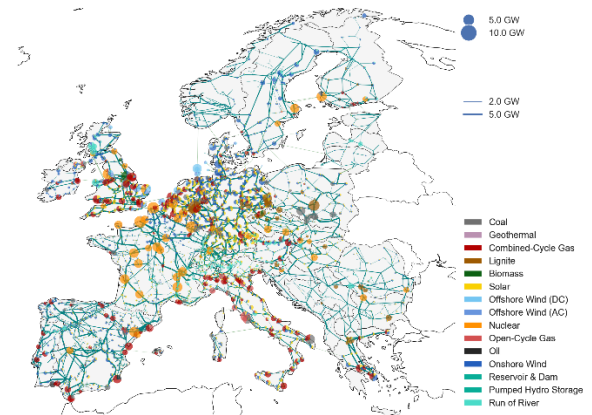


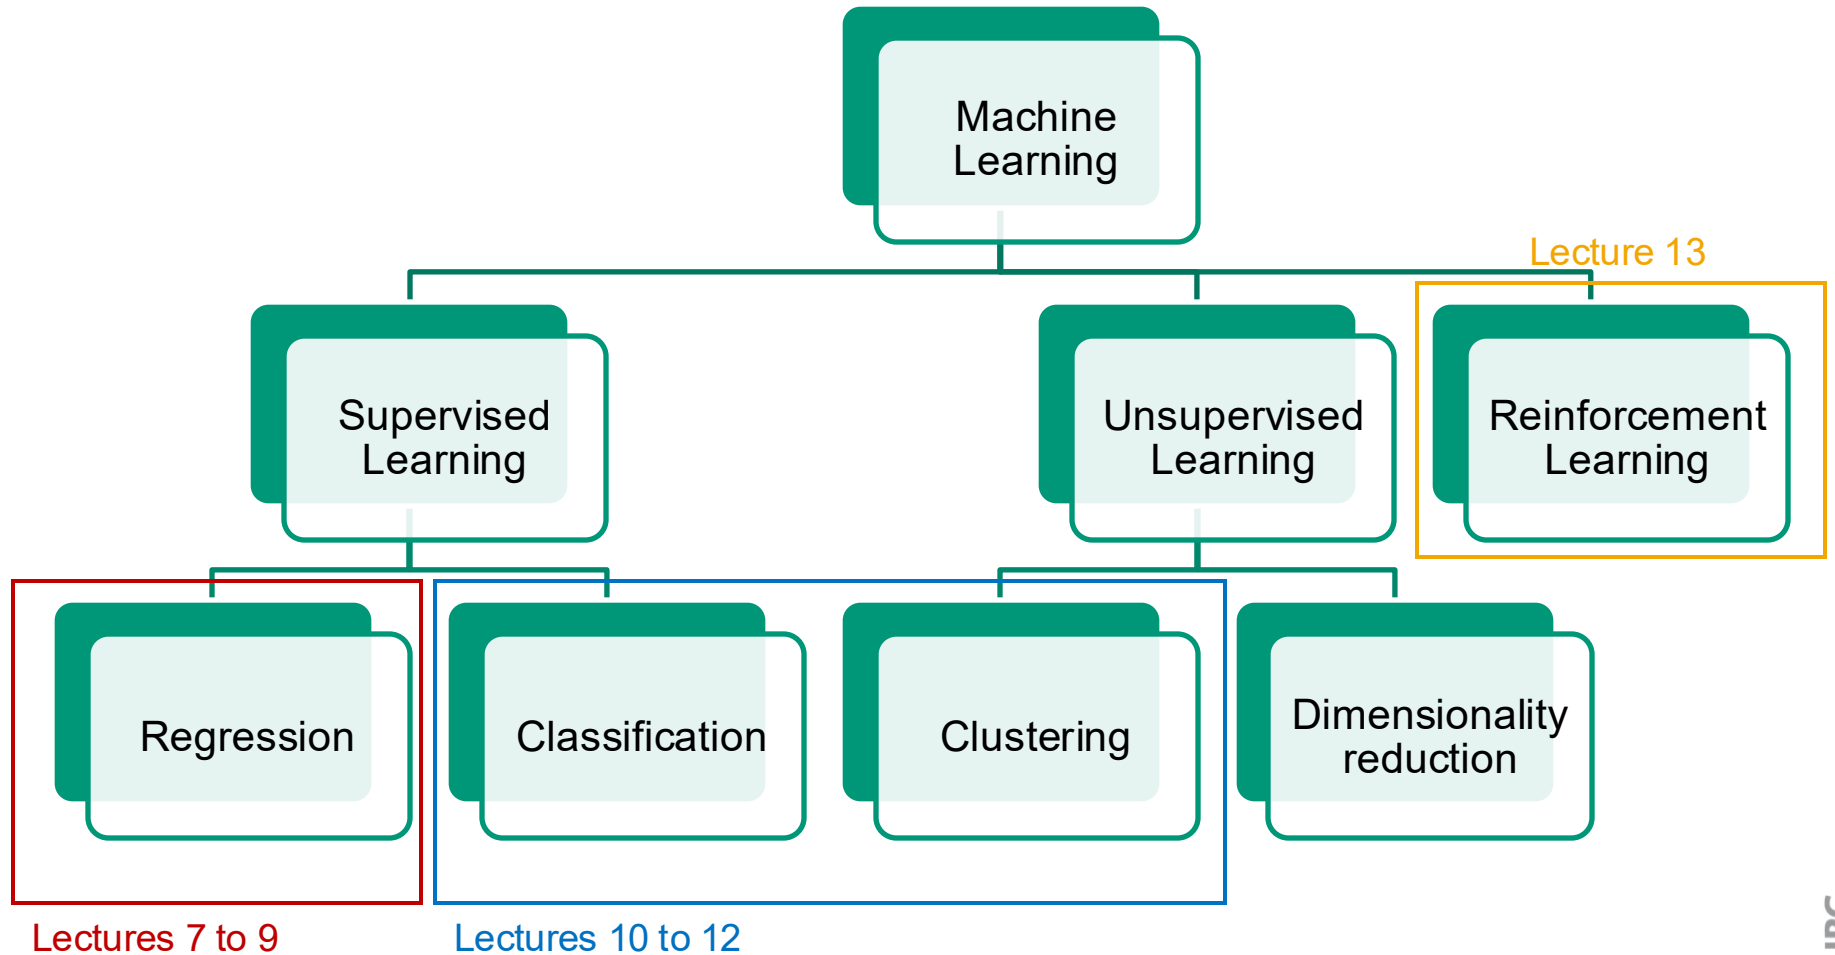
# Energy System Modeling with Python

University of Freiburg (Germany) | Faculty of Engineering  
Department of Sustainable Systems Engineering | INATECH  
**Chair for Control and Integration of Grids**

Tuesday, 3. June 2025



# Branches of ML



# Theoretical part



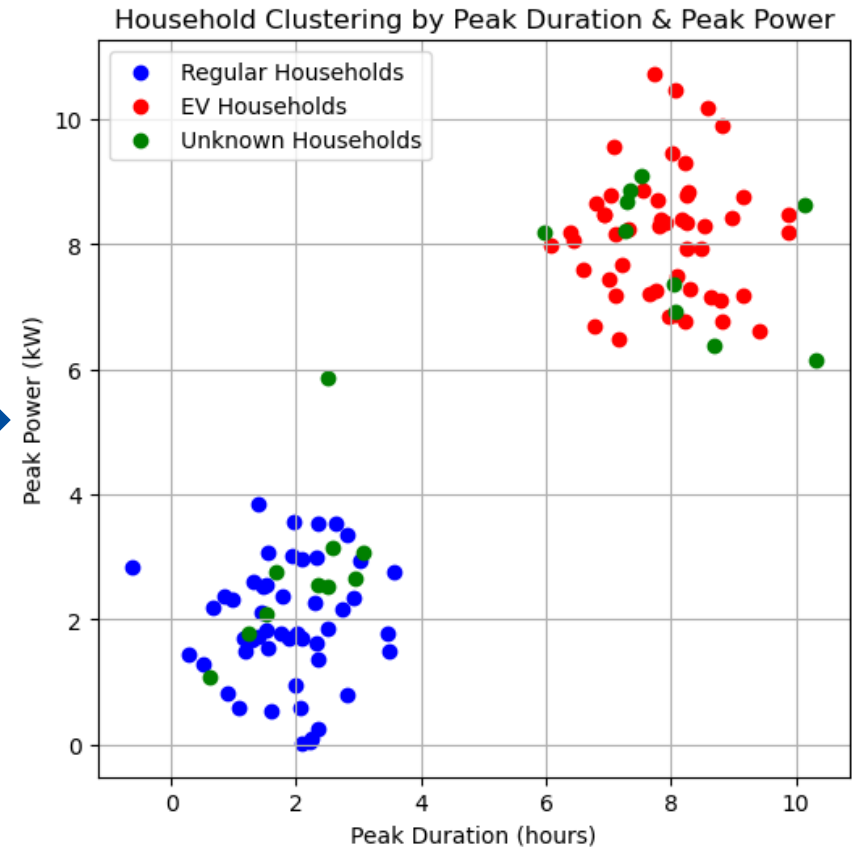

# From Un- to Supervised Learning

---

- **Last lecture:** labels were not available → **Unsupervised learning** to discover *structure* in data without labels (e.g., clustering similar load profiles)
- **Now:** labels are available → move to **supervised learning** to learn a function  $h: \mathcal{X} \rightarrow \mathcal{Y}$ , where
  - $\mathcal{X}$  = input features,
  - $\mathcal{Y}$  = discrete class labels
- **Task:** predict class  $y \in \{0, 1, \dots, K - 1\}$  from features  $x \in R^d$

Find  $h(x) \approx y$  with  $y \in \{\text{class labels}\}$

# Problem Formulation



# Regression vs Classification

---

➤ **What changes when we add labels?** → Depends on the **type** of label

➤ **Regression:**

- Target variable  $y \in R$  (continuous)
- Goal: predict precise numeric output
- Typical loss: Mean Squared Error (MSE)



➤ **Classification:**

- Target variable  $y \in \{0, 1, \dots, K - 1\}$  (discrete)
- Goal: assign input to one of K classes
- Typical loss: Cross-Entropy or 0-1 loss (more later)

➤ **Consequences for model behavior:**

- Regression → continuous output
- Classification → probability or discrete label

# Regression vs Classification



# Energy-System Examples: Regression vs Classification

Task	Label Type	Goal	Method Type
Demand Forecasting	Continuous (kWh)	Predict numeric load	Regression
Device Type Classification	Multi-class (e.g., heat pump, boiler, ...)	Identify type	Classification
Anomaly Detection	Binary or multi-class	Detect and classify faults	Classification
Price Estimation	Continuous (€/MWh)	Predict future price	Regression

- ✓ Choosing the correct ML formulation depends on the *nature of the label*
- ✓ Common tasks in energy systems involve **both** regression and classification



# Loss & Cost Refresher

---

- You've already seen this in **regression**:
  - *Loss* = error for one example
  - *Cost* = average loss over training set
- Same structure holds for classification, but with different loss functions
- Classification requires **losses suitable for discrete outputs**
- We'll look at:
  - 0-1 Loss (ideal, non-differentiable)
  - Hinge Loss (Support Vector Machine SVM)
  - **Cross-Entropy Loss** (used today)

# Classification Loss Functions – Overview

---

## 0–1 Loss

$$\mathcal{L}(y, \hat{y}) = \begin{cases} 0 & \text{if } y = \hat{y} \\ 1 & \text{otherwise} \end{cases}$$

- Directly measures classification accuracy
- Not differentiable → not usable for gradient descent

## Hinge Loss (SVMs)

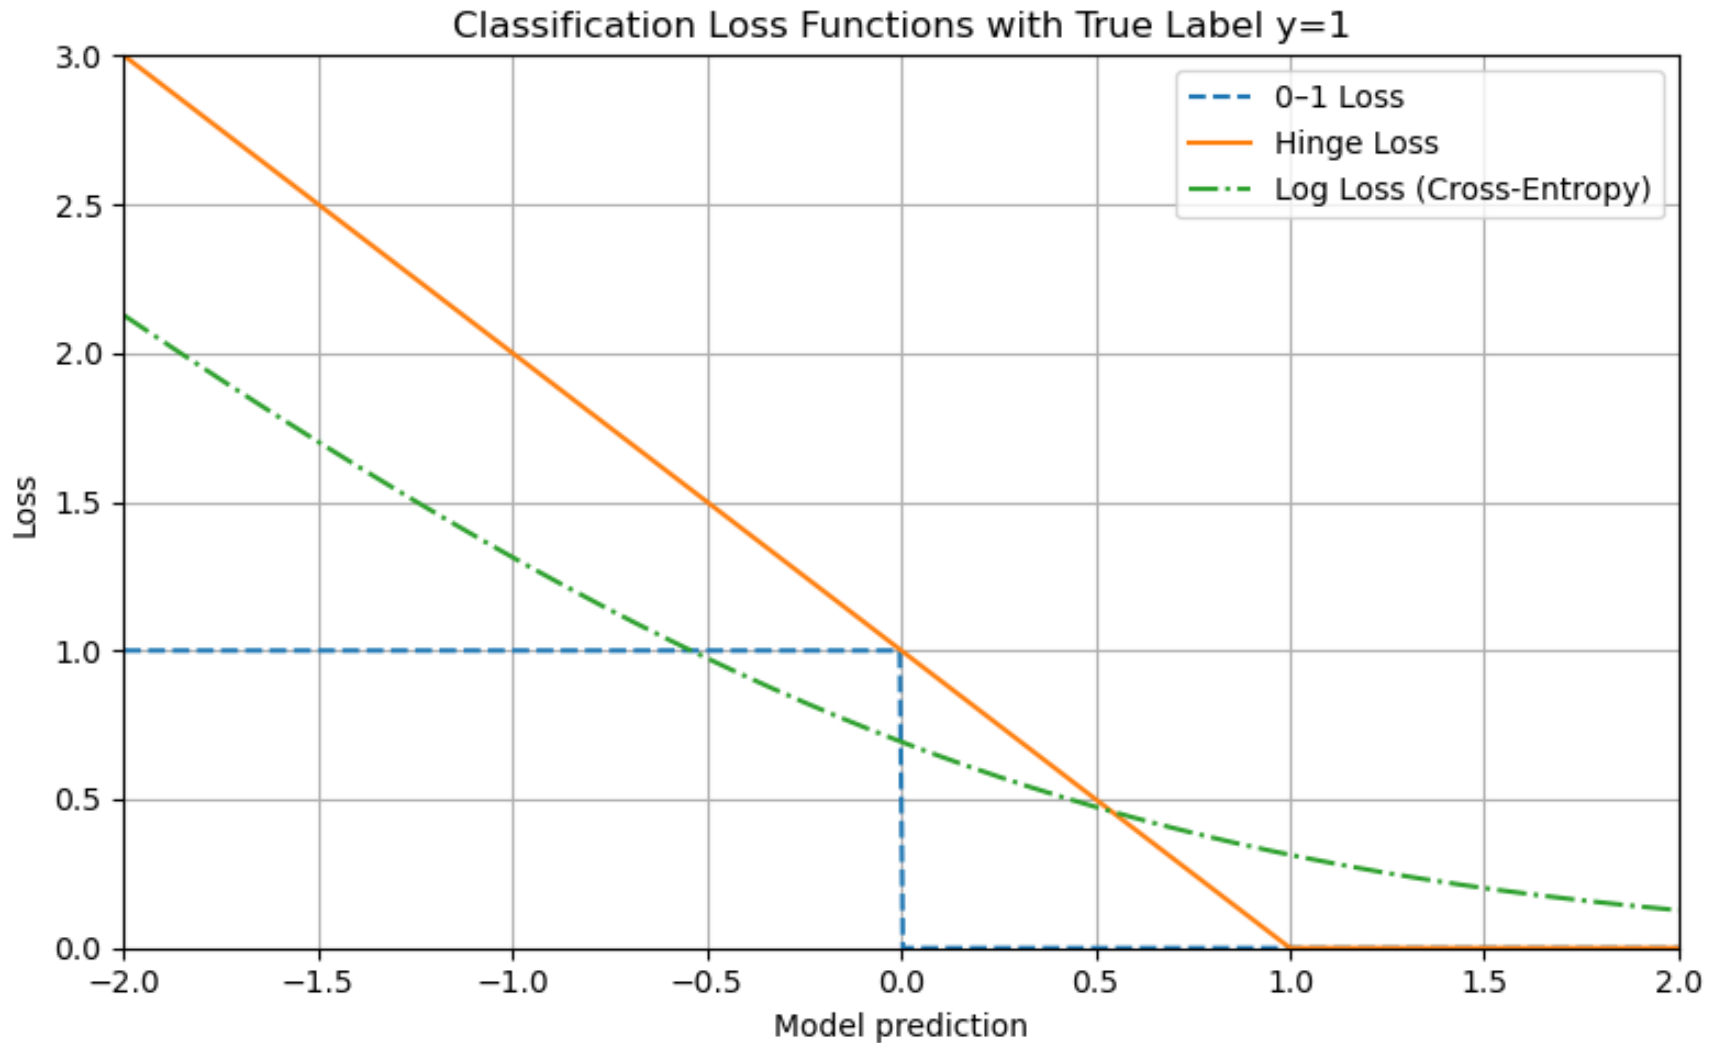
$$\mathcal{L}(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y})$$

- Promotes large margins
- Used in Support Vector Machines

## Log Loss / Cross-Entropy (our focus)

- Smooth, convex
- Measures probability error
- Used in logistic regression & neural networks

# Classification Loss Functions – Overview



# Cross-Entropy Loss (Binary Case)

---

- Used in **binary classification**, especially with **logistic regression**
- Prediction using **sigmoid function**:

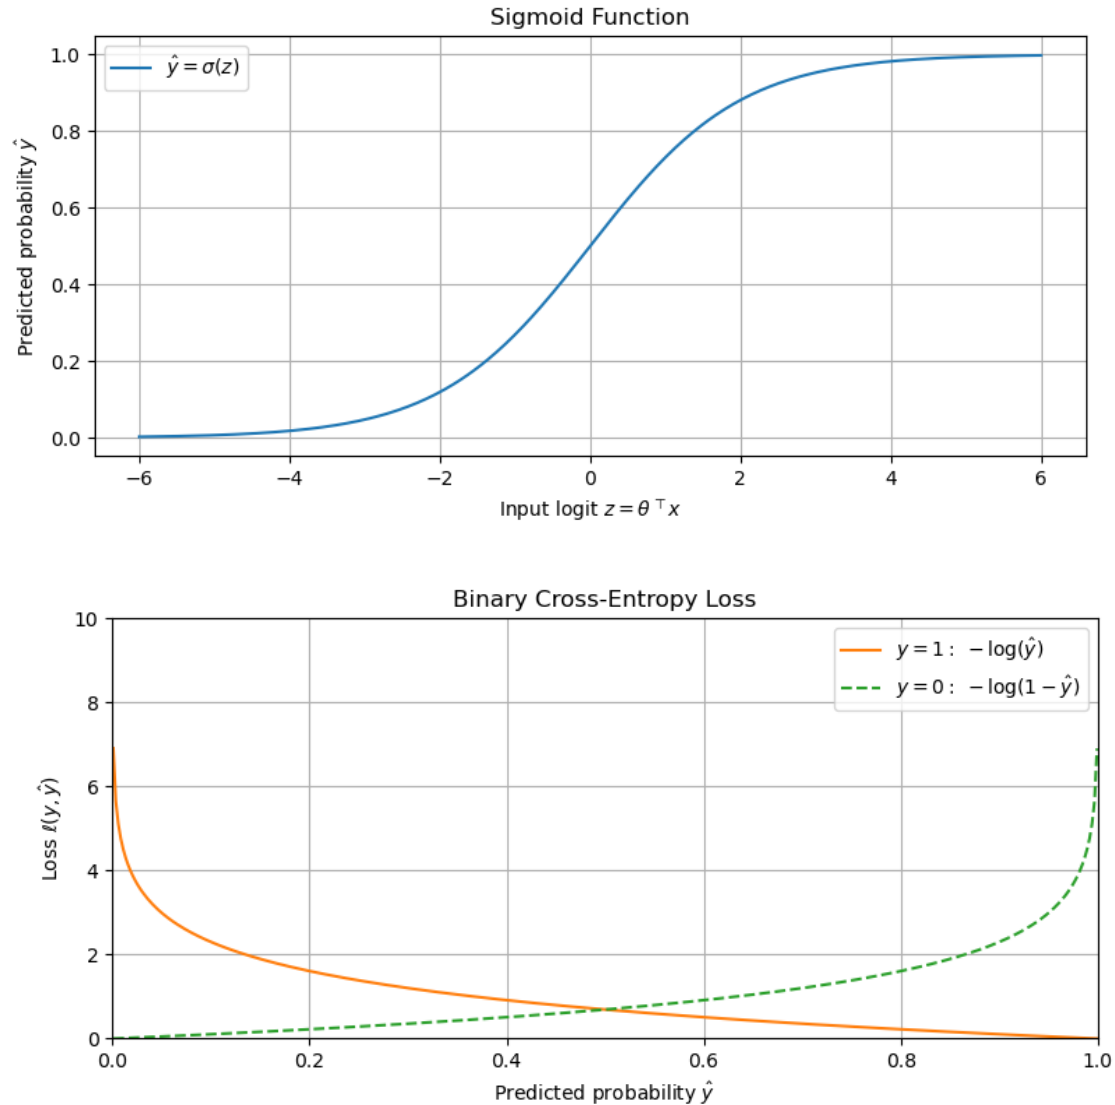
$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}, \quad z = \theta^T x$$

- **Binary Cross-Entropy Loss:**

$$\mathcal{L}(y, \hat{y}) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

- **Intuition:**
  - Penalizes confident wrong predictions heavily
  - Encourages well-calibrated probabilities

# Cross-Entropy Loss (Binary Case)



# Classification Models



# Logistic Regression – Probabilistic Classifier

- **Model form:**

$$\hat{y} = \sigma(z) = \frac{1}{1+e^{-z}} \text{ where } z = \theta^T x$$

- **Training objective:** minimize **binary cross-entropy**

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

- Optimized via **gradient-based methods** (e.g. gradient descent)

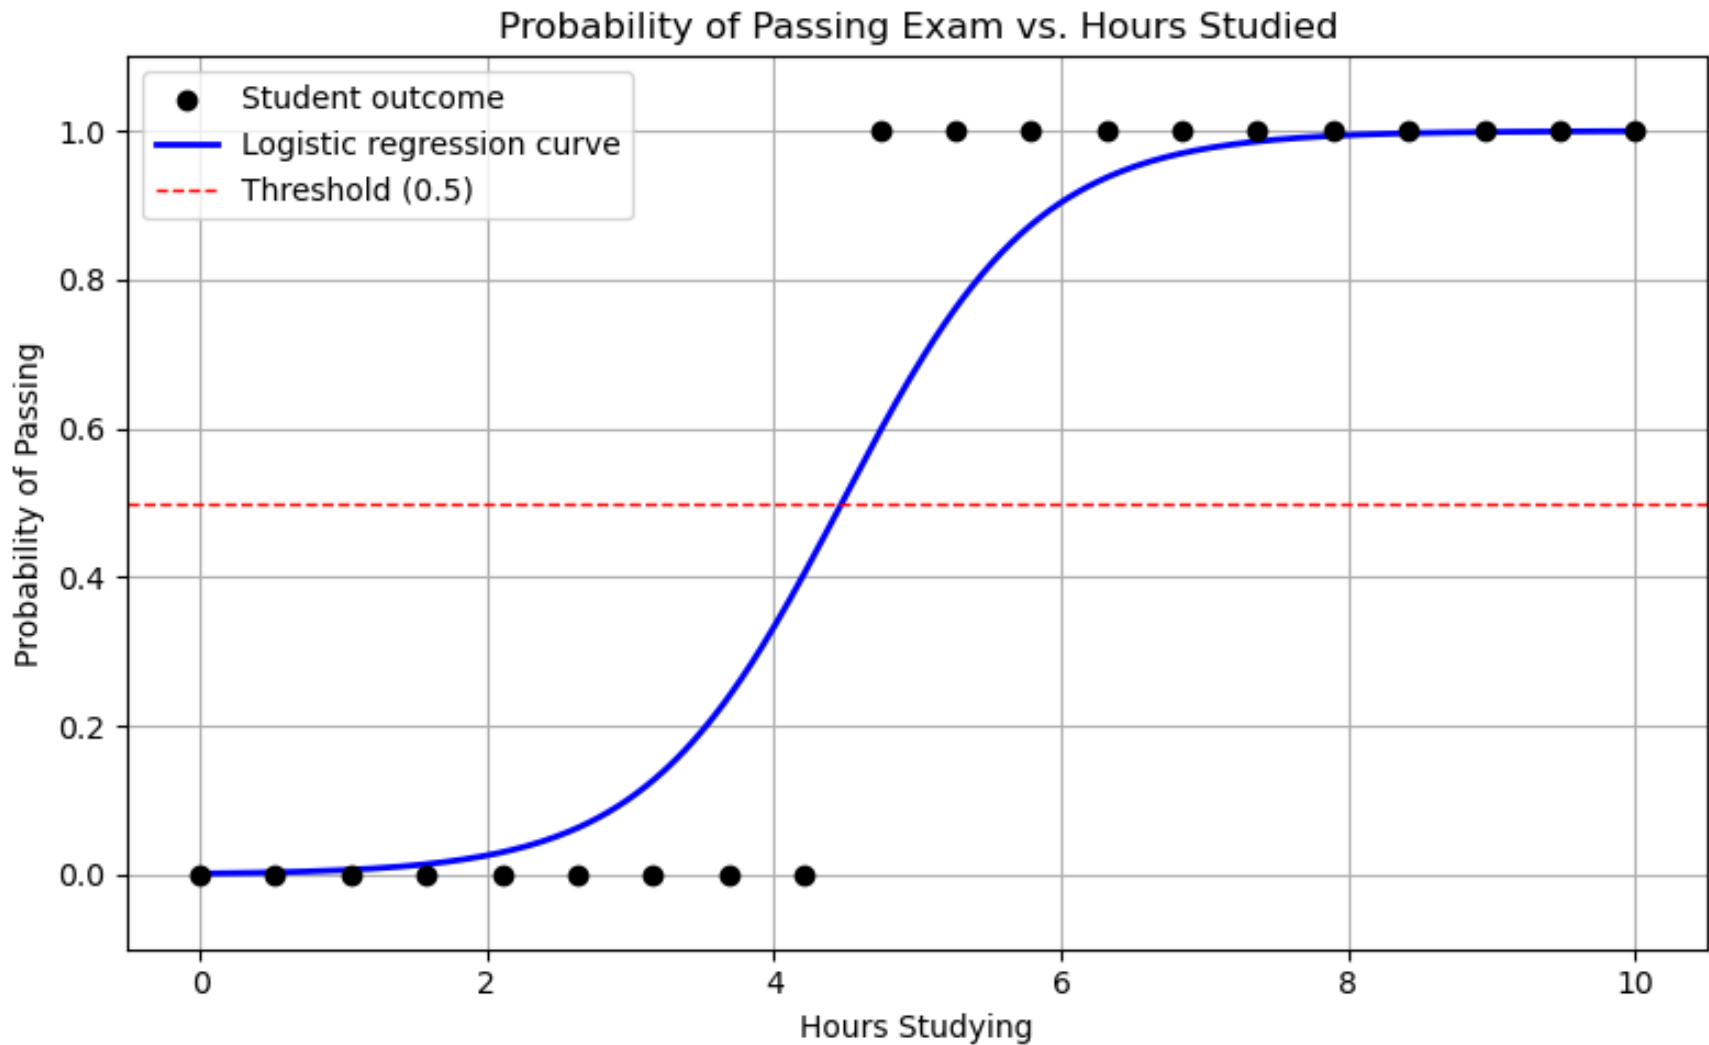
- Output  $\hat{y} \in (0,1)$  is a **probability** estimate for class 1

- **Decision rule:**

$$\text{Predict class} = \begin{cases} 1 & \text{if } \hat{y} \geq \tau \\ 0 & \text{otherwise} \end{cases}$$

- Logistic regression learns a **linear decision boundary** in feature space

# Logistic Regression – Probabilistic Classifier







# Random Forest Classifier – Splits & Impurity

Toy Tree Node: Gini Impurity Before and After Split

- **Ensemble of trees** built on bootstrapped samples & random feature subsets
- Trees choose splits that **maximize impurity reduction**
- **Node impurity** quantifies class mix:
  - **Gini:**  $1 - \sum_k p_k^2$
  - **Entropy:**  $-\sum_k p_k \log_2 p_k$
- **Regression trees** instead minimize **MSE** at each split




# Random Forest Classifier – Voting & Probabilities

- **Hard vote:** final class = majority of tree predictions

- **Soft vote:** class k probability

$$P(y = k | x) = \frac{1}{T} \sum_{t=1}^T 1 \{h_t(x) = k\}$$


- Probabilities enable thresholding and confidence analysis
- **Regression trees** instead use **average** of the trees



# LightGBM for Classification – What Changes?

- LightGBM = **gradient boosting** → models built sequentially, each new tree fits the **gradient of the loss**
- You've seen this for regression:
  - Loss: **MSE**
  - Gradients: residuals  $y_i - \hat{y}_i$
- For **classification**:
  - Loss = **binary cross-entropy**  
$$\mathcal{L}(y, \hat{y}) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$
  - Gradients = from this loss, based on predicted probabilities
- Output is a **logit score**, converted via sigmoid:
$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$
- Supports:
  - **Hard predictions** via thresholding (e.g., 0.5)
  - **Probabilistic outputs** for metric evaluation

# Performance Metrics



# Confusion Matrix Fundamentals

- **Definition of terms**
  - **True Positive (TP):** model predicts 1, actual = 1
  - **False Positive (FP):** model predicts 1, actual = 0
  - **False Negative (FN):** model predicts 0, actual = 1
  - **True Negative (TN):** model predicts 0, actual = 0
- **Purpose:** all classification metrics derive from these four counts
- **Cost implications:** different errors carry different real-world costs (e.g., missed fault vs. false alarm)

	Predicted 1	Predicted 0
Actual 1	TP	FN
Actual 0	FP	TN

# Precision & Recall Trade-off

- **Precision (P):** also known as positive predictive value

$$P = \frac{TP}{TP + FP}$$

- ✓ “Of all predicted positives, how many were correct?”

- **Recall (R):** also known as sensitivity

$$R = \frac{TP}{TP + FN}$$

- ✓ “Of all actual positives, how many did we catch?”

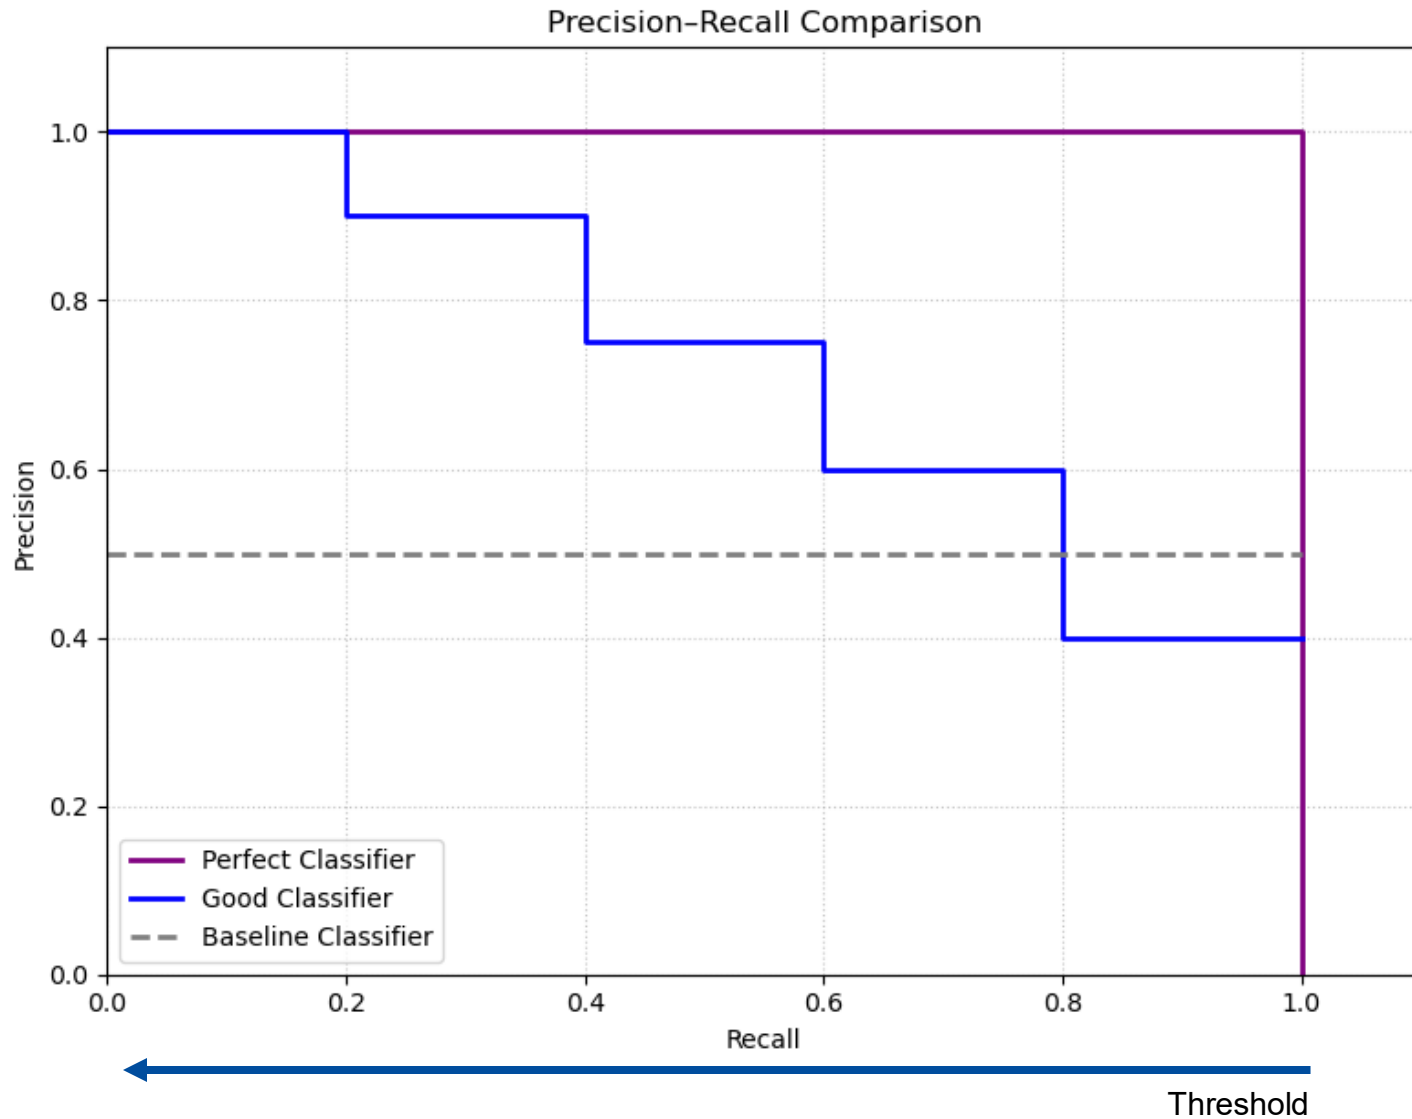
- **Trade-offs:**

- ✓ raising threshold → ↑precision, ↓recall;
- ✓ lowering threshold → ↑recall, ↓precision

- **When to prioritize:**

- ✓ High precision → costly false positives (e.g., dispatching crews)
- ✓ High recall → costly misses (e.g., undetected faults)

# Precision & Recall Trade-off



# F1 Score & Averaging

- **F1 score** combines precision and recall into one metric:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

- **Why F1?** Balances false positives and false negatives when both matter.

## Averaging for multi-class:

- **Macro-F1**

- Compute F1 for each class independently and take the **unweighted mean**:

$$\text{Macro-F1} = \frac{1}{K} \sum_{k=1}^K F1_k$$

- Treats all classes equally → highlights poor performance on rare classes

- **Weighted-F1**

- Compute F1 for each class, then weight by support  $n_k$ :

$$\text{Weighted-F1} = \frac{1}{N} \sum_{k=1}^K n_k F1_k, \text{ where } N = \sum_k n_k$$

- Reflects dataset composition → gives more influence to common classes



# F1 Score & Averaging

## Averaging for multi-class:

### ➤ Macro-F1

- Compute F1 for each class independently and take the **unweighted mean**:

$$\text{Macro-F1} = \frac{1}{K} \sum_{k=1}^K F1_k$$

- Treats all classes equally → highlights poor performance on rare classes

### ➤ Weighted-F1

- Compute F1 for each class, then weight by support  $n_k$ :

$$\text{Weighted-F1} = \frac{1}{N} \sum_{k=1}^K n_k F1_k, \text{ where } N = \sum_k n_k$$

- Reflects dataset composition → gives more influence to common classes

Class	F1	Support $n_k$
A	0.90	100
B	0.60	20
Macro-F1	$(0.90+0.60)/2 = 0.75$	—
Weighted-F1	$(100 \cdot 0.90 + 20 \cdot 0.60)/120 = 0.85$	—

# Example – Confusion Matrix & Classification Report

## Confusion Matrix

Predicted ↓ \ True →	Baseline	EV
Baseline	191	9
EV	9	41

## Classification Report

Class	Precision	Recall	F1-Score	Support
Baseline	0.95	0.95	0.95	200
EV	0.82	0.82	0.82	50
<b>Accuracy</b>			<b>0.93</b>	250
<b>Macro avg</b>	0.89	0.89	0.89	250
<b>Weighted avg</b>	0.93	0.93	0.93	250

## Key Takeaways from Classification Report

- **High overall accuracy (0.93) hides class imbalance effects**
- **Baseline class (n=200):** precision & recall = 0.95 → very reliable
- **EV class (n=50):** precision & recall = 0.82 → notably weaker performance
- **Macro-F1 = 0.89** < overall accuracy → equal weighting reveals minority-class challenges
- **Weighted-F1 = 0.93** ≈ accuracy → majority class dominates the aggregate metric
- **Next steps:** focus on improving EV detection (e.g., resampling, class weights, tailored features)

# Unbalanced Datasets



# Imbalanced Datasets – Why It's a Problem

Confusion Matrix

Predicted ↓ \ True →	Baseline	EV
Baseline	90	10
EV	0	0

Classification Report

Class	Precision	Recall	F1-Score	Support
Baseline	0.90	1.0	0.95	90
EV	0.00	0.00	0.00	10
<b>Accuracy</b>			<b>0.90</b>	100
<b>Macro avg</b>	0.45	0.50	0.47	100
<b>Weighted avg</b>	0.81	0.90	0.85	100

- **Definition:** one class (majority) far outnumber(s) the other(s) (minority)
- **Naïve accuracy trap:**
  - Suppose Baseline:EV = 90:10 ratio
  - Always predicting “Baseline” → 90% accuracy, but **zero** detection of EV
- **Real-world stakes:**
  - In **fault detection**, rare faults may be only 1–5% of data
  - Missing every fault → catastrophic consequences despite high accuracy
- **Key insight:** accuracy alone is **misleading** under imbalance

# Coffee Break

---



# Time to put everything into code

**START  
CODING  
SESSION**



# What You'll Do in Code Today

---

## Goal:

Assign missing households their correct demand-response device type.

## Main Steps:

1. **Load data** (household profiles + known device labels)
2. **Feature extraction** (reuse last exercise's time-series summary pipeline)
3. **Split** into training and test sets (stratified by device type)
4. **Train** three classifiers:
  - a. **Logistic Regression**
  - b. **Random Forest Classifier**
  - c. **LightGBM Classifier**
5. **Evaluate** on test set:
  - a. **Confusion matrix** for each model
  - b. **Classification report** (precision, recall, F1, support)
  - c. **Macro-F1** and **weighted-F1** comparisons
6. **Compare** model behavior and decision making

# Takeaways

---

- ✓ Classification tasks require **appropriate loss functions** (cross-entropy) and **evaluation metrics** beyond accuracy
- ✓ **Logistic Regression, Random Forest, and LightGBM** offer different trade-offs in interpretability, non-linearity, and handling of rare classes
- ✓ **Macro-F1 vs Weighted-F1** reveal performance on minority classes
- ✓ **Unbalanced datasets** can often be an issue and require special treatment

## Further Questions to Think About

- How might you **improve detection** of the rarest device class?
- When would you choose **one classifier** over another in production?
- How could **threshold tuning** or **class weighting** further boost performance?
- What additional **features** might capture device-specific behavior more effectively?